

Chapter 06

(version 11th October 2001)

Maarten van Steen

Vrije Universiteit Amsterdam, Faculty of Science

Dept. Mathematics and Computer Science

Room R4.20. Tel: (020) 444 7784

E-mail: steen@cs.vu.nl, URL: www.cs.vu.nl/~steen

01	Introduction
02	Communication
03	Processes
04	Naming
05	Synchronization
06	Consistency and Replication
07	Fault Tolerance
08	Security
09	Distributed Object-Based Systems
10	Distributed File Systems
11	Distributed Document-Based Systems
12	Distributed Coordination-Based Systems

00 – 1

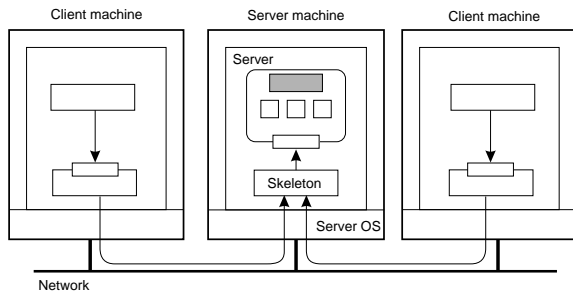
/

Consistency & Replication

- Introduction (what's it all about)
- Data-centric consistency
- Client-centric consistency
- Distribution protocols
- Consistency protocols
- Examples

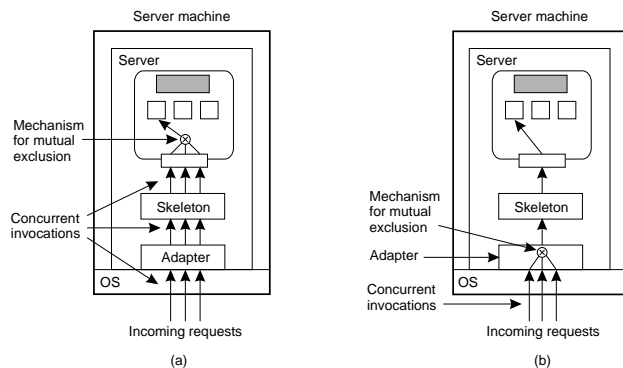
Shared Objects

Problem: If objects (or data) are shared, we need to do something about concurrent accesses to guarantee state consistency.




Concurrency Control (1/2)

Problem: Is the remote object already thread-safe or not?



Concurrency Control (2/2)

Problem: Should we seek for object-specific solutions, or generally applicable ones?



Question: Why would we want object-specific replication protocols?

Performance and Scalability

Main issue: To keep replicas consistent, we generally need to ensure that all *conflicting* operations are done in the the same order everywhere

Conflicting operations: From the world of transactions:

- Read–write conflict: a read operation and a write operation act concurrently
- Write–write conflicts: two concurrent write operations

Guaranteeing global ordering on conflicting operations may be a costly operation, downgrading scalability


Solution: weaken consistency requirements so that hopefully global synchronization can be avoided

Data-Centric Consistency Models

(1/2)

Consistency model: a contract between a (distributed) data store and processes, in which the data store specifies precisely what the results of read and write operations are in the presence of concurrency.

Essence: A data store is a distributed collection of storages accessible to clients:



Data-Centric Consistency Models

(2/2)

Strong consistency models: Operations on shared data are synchronized:

- Strict consistency (related to time)
- Sequential consistency (what we are used to)
- Causal consistency (maintains only causal relations)
- FIFO consistency (maintains only individual ordering)

Weak consistency models: Synchronization occurs only when shared data is locked and unlocked:


- General weak consistency
- Release consistency
- Entry consistency

Observation: The weaker the consistency model, the easier it is to build a scalable solution.

Strict Consistency

Any read to a shared data item X returns the value stored by the most recent write operation on X .

Observation: It doesn't make sense to talk about "the most recent" in a distributed environment.




- Assume all data items have been initialized to NIL
- $W(x)a$: value a is written to x
- $R(x)a$: reading x returns the value a

Note: Strict consistency is what you get in the normal sequential case, where your program does not interfere with any other program.

Sequential Consistency

The result of any execution is the same as if the operations of all processes were executed in some sequential order, and the operations of each individual process appear in this sequence in the order specified by its program.



Note: We're talking about **interleaved** executions: there is some total ordering for all operations taken together.



Linearizable: Sequential plus operations are ordered according to a global time.


Causal Consistency

Writes that are potentially causally related must be seen by all processes in the same order. Concurrent writes may be seen in a different order by different processes.



FIFO Consistency

Writes done by a single process are received by all other processes in the order in which they were issued, but writes from different processes may be seen in a different order by different processes.



Weak Consistency (1/2)

- Accesses to **synchronization variables** are *sequentially consistent*.
- No access to a synchronization variable is allowed to be performed until all previous writes have completed everywhere.
- No data access is allowed to be performed until all previous accesses to synchronization variables have been performed.

Basic idea: You don't care that reads and writes of a *series* of operations are immediately known to other processes. You just want the *effect* of the series itself to be known.

Weak Consistency (2/2)

P1:	W(x)a	W(x)b	S		
P2:			R(x)a	R(x)b	S
P3:			R(x)b	R(x)a	S

(a)

P1:	W(x)a	W(x)b	S		
P2:			S	R(x)a	

(b)

Observation: Weak consistency implies that we need to lock and unlock data (implicitly or not).

Release Consistency

Idea: Divide access to a synchronization variable into two parts: an **acquire** and a **release** phase. Acquire forces a requester to wait until the shared data can be accessed; release sends requester's local value to other servers in data store.

P1:	Acq(L)	W(x)a	W(x)b	Rel(L)
P2:			Acq(L)	R(x)b
P3:				Rel(L)
				R(x)a

Entry Consistency

- With release consistency, all local updates are propagated to other copies/servers during release of shared data.
- With entry consistency, each shared data item is associated with a synchronization variable.
- When acquiring the synchronization variable, the most recent values of its associated shared data item are fetched.

Note: Where release consistency affects *all* shared data, entry consistency affects only those shared data associated with a synchronization variable.

Question: What would be a convenient way of making entry consistency more or less transparent to programmers?

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

[illegible]

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.


- [illegible]

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

[illegible]


This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

Basic Architecture



Monotonic Reads (1/2)

If a process reads the value of a data item x , any successive read operation on x by that process will always return that same or a more recent value.




Notation: $WS(x_i[t])$ is the set of write operations (at L_i) that lead to version x_i of x (at time t); $WS(x_i[t_1];x_j[t_2])$ indicates that it is known that $WS(x_i[t_1])$ is part of $WS(x_j[t_2])$.

Note: Parameter t is omitted from figures

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.



[illegible][illegible][illegible]

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

[illegible]

Read Your Writes



The effect of a write operation by a process on data item x , will always be seen by a successive read operation on x by the same process.



Example: Updating your Web page and guaranteeing that your Web browser shows the newest version instead of its cached copy.

Writes Follow Reads

A write operation by a process on a data item x following a previous read operation on x by the same process, is guaranteed to take place on the same or a more recent value of x that was read.



Example: See reactions to posted articles only if you have the original posting (a read “pulls in” the corresponding write operation).

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

-
- This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.


This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.


This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

-
- This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.



This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.



This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

-
- This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

06 – 30 Consistency & Replication/6.4 Distribution Protocols

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

- 06 – 31 Consistency & Replication/6.4 Distribution Protocols

Issue	Push-based	Pull-based
1:	List of client caches	None
2:	Update (and possibly fetch update)	Poll and update
3:	Immediate (or fetch-update time)	Fetch-update time

1: *State at server*
2: *Messages to be exchanged*
3: *Response time at the client*

Update Propagation (3/3)

Observation: We can dynamically switch between pulling and pushing using **leases**: A contract in which the server promises to push updates to the client until the lease expires.

Issue: Make lease expiration time dependent on system's behavior (adaptive leases):

- Age-based leases: An object that hasn't changed for a long time, will not change in the near future, so provide a long-lasting lease
- Renewal-frequency based leases: The more often a client requests a specific object, the longer the expiration time for that client (for that object) will be
- State-based leases: The more loaded a server is, the shorter the expiration times become

Question: Why are we doing all this?

Epidemic Algorithms

- General background
- Update models
- Removing objects

[illegible]

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

- [illegible]

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

[illegible]

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

-
- This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

Anti-Entropy

Basic issue: When a server S contacts another server S^* to exchange state information, three different strategies can be followed:

Push: S only forwards all its updates to S^* :

if $T(O, S^*) < T(O, S)$
then $VAL(O, S^*) \leftarrow VAL(O, S)$

Pull: S only fetches updates from S^* :

if $T(O, S^*) > T(O, S)$
then $VAL(O, S) \leftarrow VAL(O, S^*)$

Push-Pull: S and S^* exchange their updates by pushing and pulling values

Observation: if each server periodically randomly chooses another server for exchanging updates, an update is propagated in $O(\log(N))$ time units.

Question: Why is pushing alone not efficient when many servers have already been updated?

Gossiping

Basic model: A server S having an update to report, contacts other servers. If a server is contacted to which the update has already propagated, S stops contacting other servers with probability $1/k$.

If s is the fraction of ignorant servers (i.e., which are unaware of the update), it can be shown that with many servers:

$$s = e^{-(k+1)(1-s)}$$

k	s
1	0.2000
2	0.0600
3	0.0200
4	0.0070
5	0.0025

Observation: If we really have to ensure that all servers are eventually updated, gossiping alone is not enough

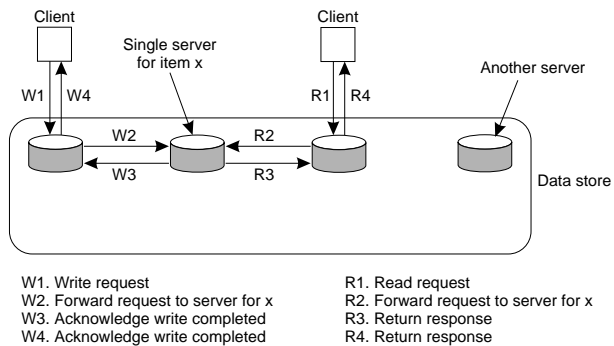
This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

This image shows a single sheet of white paper with horizontal blue or grey ruling lines. The lines are evenly spaced and run across the width of the page. There is no handwriting or other markings on the paper.

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

Primary-Based Protocols (1/4)

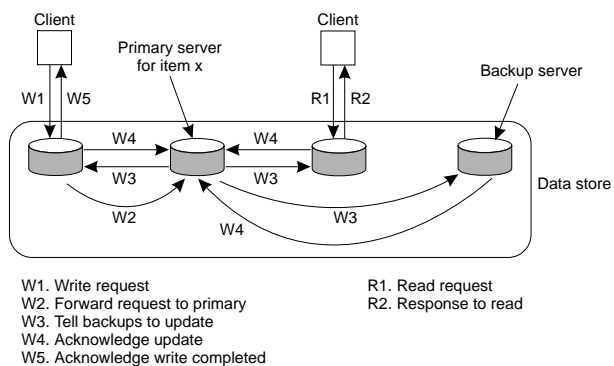
Primary-based, remote-write, fixed server:



Example: Used in traditional client-server systems that do not support replication.

Primary-Based Protocols (2/4)


Primary-backup protocol:



Example: Traditionally applied in distributed databases and file systems that require a high degree of fault tolerance. Replicas are often placed on same LAN.

Primary-Based Protocols (3/4)


Primary-based, local-write protocol:



Example: Establishes only a fully distributed, non-replicated data store. Useful when writes are expected to come in series from the same client (e.g., mobile computing without replication)

Primary-Based Protocols (4/4)


Primary-backup protocol with local writes:



Example: Distributed shared memory systems, but also mobile computing in disconnected mode (ship all relevant files to user before disconnecting, and update later on).

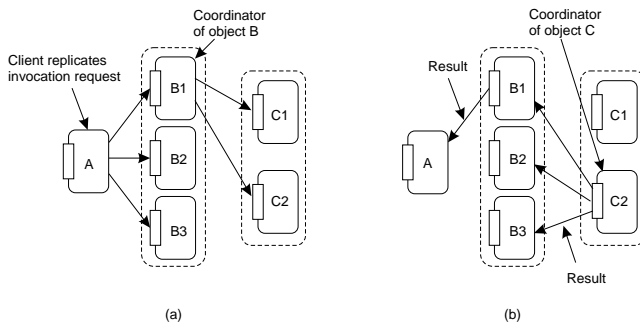
Replicated-Write Protocols (1/3)

Active replication: Updates are forwarded to multiple replicas, where they are carried out. There are some problems to deal with in the face of replicated invocations:




Replicated-Write Protocols (2/3)

Replicated invocations: Assign a coordinator on each side (client and server), which ensures that only one invocation, and one reply is sent:




Replicated-Write Protocols (3/3)

Quorum-based protocols: Ensure that each operation is carried out in such a way that a majority vote is established: distinguish **read quorum** and **write quorum**:



Example: Lazy Replication

Basic model: Number of replica servers jointly implement a causal-consistent data store. Clients normally talk to **front ends** which maintain data to ensure causal consistency.



Lazy Replication: Vector Timestamps

VAL(i): VAL(i)[i] denotes the total number of write operations sent directly by a front end (client). VAL(i)[j] denotes the number of updates sent from replica #j.


WORK(i): WORK(i)[i] total number of write operations directly from front ends, including the pending ones. WORK(i)[j] is total number of updates from replica #j, including pending ones.

LOCAL(C): LOCAL(C)[j] is (almost) most recent value of VAL(j)[j] known to front end C (will be refined in just a moment)

DEP(R): Timestamp associated with a request, reflecting what the request depends on.

Operations

Read operations:



Write operations:

