

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

(version 27th November 2001)

This image shows a single sheet of white paper with horizontal blue or grey ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

Dept. Mathematics and Computer Science

E-mail: steen@cs.vu.nl, URL: www.cs.vu.nl/~steen/

- [illegible]

/

This image shows a single sheet of white paper with horizontal blue or grey ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

-
- This image shows a single sheet of white paper with horizontal blue or grey ruling lines. The lines are evenly spaced and run across the width of the page. There are approximately 20 lines visible. The paper has a slight shadow on its right side, suggesting it's resting on a surface.

CORBA


CORBA: Common **O**bject **R**equest **B**roker **A**rchitecture

Background:

- Developed by the **Object Management Group (OMG)** in response to industrial demands for object-based middleware
- Currently in version #2.4 with #3 (almost) done
- CORBA is a *specification*: different implementations of CORBA exist
- Very much the work of a committee: there are over 800 members of the OMG and many of them have a say in what CORBA should look like

Essence: CORBA provides a simple distributed-object model, with specifications for many supporting services ⇒ it may be here to stay (for a couple of years)

CORBA Overview (1/2)



Object Request Broker (ORB): CORBA's object broker that connects clients, objects, and services

Proxy/Skeleton: Precompiled code that takes care of (un)marshaling invocations and results

Dynamic Invocation/Skeleton Interface (DII/DSI): To allow clients to "construct" invocation requests at runtime instead of calling methods at a proxy, and having the server-side "reconstruct" those request into regular method invocations

Object adapter: Server-side code that handles incoming invocation requests.

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

CORBA Services

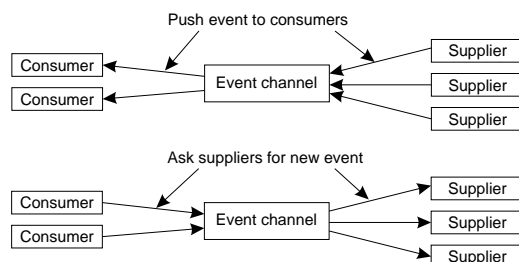
| Service | Description |
|-----------------|--|
| Collection | Facilities for grouping objects into lists, queue, sets, etc. |
| Query | Facilities for querying collections of objects in a declarative manner |
| Concurrency | Facilities to allow concurrent access to shared objects |
| Transaction | Flat and nested transactions on method calls over multiple objects |
| Event | Facilities for asynchronous communication through events |
| Notification | Advanced facilities for event-based asynchronous communication |
| Externalization | Facilities for marshaling and unmarshaling of objects |
| Life cycle | Facilities for creation, deletion, copying, and moving of objects |
| Licensing | Facilities for attaching a license to an object |
| Naming | Facilities for systemwide naming of objects |
| Property | Facilities for associating (attribute, value) pairs with objects |
| Trading | Facilities to publish and find the services an object has to offer |
| Persistence | Facilities for persistently storing objects |
| Relationship | Facilities for expressing relationships between objects |
| Security | Mechanisms for secure channels, authorization, and auditing |
| Time | Provides the current time within specified error margins |

Communication Models (1/2)

Object invocations: CORBA distinguishes three different forms of direct invocations:



| Request type | Failure sem. | Description |
|----------------------|--------------|--|
| Synchronous | At-most-once | Caller blocks |
| One-way | Unreliable | Nonblocking call |
| Deferred synchronous | At-most-once | Nonblocking, but can pick-up results later |

Event communication: There are also additional facilities by means of **event channels**:



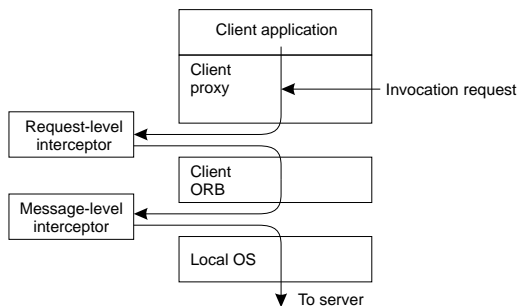
Communication Models (2/2)

Messaging facilities: reliable asynchronous and persistent method invocations:



Processes

Most aspects of processes for in CORBA have been discussed in previous classes. What remains is the concept of **interceptors**:



Request-level: Allows you to modify invocation semantics (e.g., multicasting)

Message-level: Allows you to control message-passing between client and server (e.g., handle reliability and fragmentation)

Naming

Important: In CORBA, it is essential to distinguish specification-level and implementation-level object references

Specification level: An object reference is considered to be the same as a proxy for the referenced object \Rightarrow having an object reference means you can directly invoke methods; there is no separate client-to-object binding phase


Implementation level: When a client gets an object reference, the implementation ensures that, one way or the other, a proxy for the referenced object is placed in the client's address space:

```
ObjectReference objRef;  
objRef = bindTo(object O in server S at host H);
```

Conclusion: Object references in CORBA used to be highly **implementation dependent**: different implementations of CORBA could normally not exchange their references.

Interoperable Object References (1/2)

Observation: Recognizing that object references are implementation dependent, we need a separate referencing mechanism to cross ORB boundaries

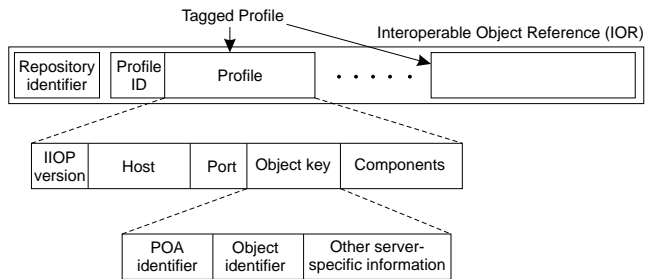


Solution: Object references passed from one ORB to another are transformed by the bridge through which they pass (different transformation schemes can be implemented)

Observation: Passing an object reference refA from ORB A to ORB B circumventing the A-to-B bridge may be useless if ORB B doesn't understand refA

Interoperable Object References (2/2)

Observation: To allow all kinds of *different* systems to communicate, we standardize the reference that is passed between bridges:



Naming Service

Essence: CORBA's naming service allows servers to associate a name to an object reference, and have clients subsequently bind to that object by resolving its name

Observation: In most CORBA implementations, object references denote servers at specific hosts; naming makes it easier to relocate objects


Observation: In the naming graph all nodes are objects; there are no restrictions to binding names to objects ⇒ CORBA allows arbitrary naming graphs

Question: How do you imagine cyclic name resolution stops?

Observation: There is no single root; an initial context node is returned through a special call to the ORB. Also: the naming service can operate *across* different ORBs ⇒ **interoperable naming service**

Fault Tolerance

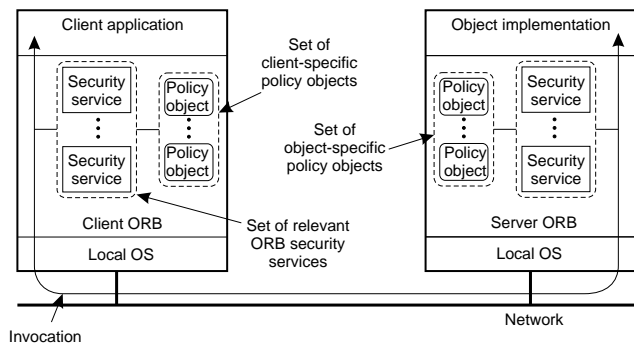
Essence: Mask failures through replication, by putting objects into **object groups**. Object groups are transparent to clients: they appear as “normal” objects. This approach requires a separate type of object reference: **Interoperable Object Group Reference:**



Note: IOGRs have the same structure as IORs; the main difference is that they are *used* differently. In IORs an additional profile is used as an alternative; in IOGR, it denotes another replica.

Security

Essence: Allow the client and object to be mostly unaware of all the security policies, except perhaps at binding time; the ORB does the rest. Specific policies are passed to the ORB as (local) objects and are invoked when necessary:



Examples: Type of message protection, lists of trusted parties.

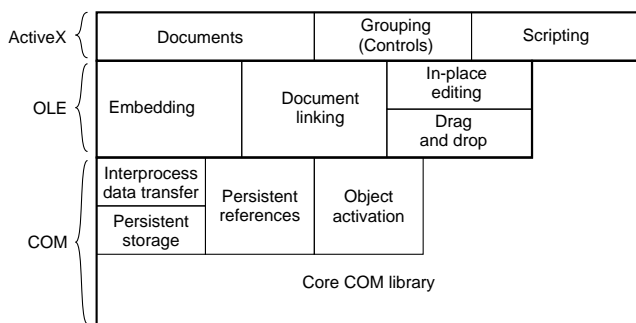
Distributed COM

DCOM: Distributed Component Object Model

- Microsoft's solution to establishing inter-process communication, possibly across machine boundaries.
- Supports a primitive notion of distributed objects
- Evolved from early Windows versions to current NT-based systems (including Windows 2000)
- Comparable to CORBA's object request broker

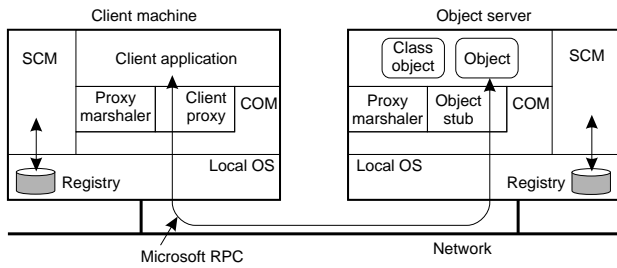
DCOM Overview (1/2)

Somewhat confused? DCOM is related to many things that have been introduced by Microsoft in the past couple of years:



DCOM: Adds facilities to communicate across process and machine boundaries.

DCOM Overview (2/2)



SCM: Service Control Manager, responsible for activating objects (cf., to CORBA's implementation repository).

Proxy marshaler: handles the way that object references are passed between different machines

COM Object Model

- An interface is a collection of semantically related operations
- Each interface is typed, and therefore has a globally unique **interface identifier**
- A client always requests an implementation of an interface:
 - Locate a class that implements the interface
 - Instantiate that class, i.e., create an object
 - Throw the object away when the client is done

DCOM Services

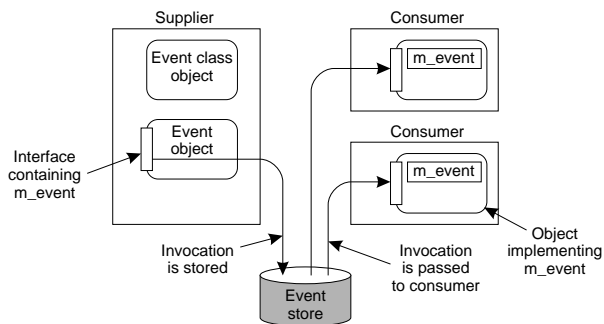
| CORBA | DCOM/COM+ | Windows 2000 |
|-----------------|---------------------------------|-------------------------------------|
| Collection | ActiveX Data Objects | – |
| Query | None | – |
| Concurrency | Thread concurrency | – |
| Transaction | COM+ Automatic Transactions | Distributed Transaction Coordinator |
| Event | COM+ Events | – |
| Notification | COM+ Events | – |
| Externalization | Marshaling utilities | – |
| Life cycle | Class factories, JIT activation | – |
| Licensing | Special class factories | – |
| Naming | Monikers | Active Directory |
| Property | None | Active Directory |
| Trading | None | Active Directory |
| Persistence | Structured storage | Database access |
| Relationship | None | Database access |
| Security | Authorization | SSL, Kerberos |
| Time | None | None |

Note: COM+ is effectively COM plus services that were previously available in an ad-hoc fashion

Communication Models

Object invocations: Synchronous remote-method calls with at-most-once semantics. Asynchronous invocations are supported through a polling model, as in CORBA.

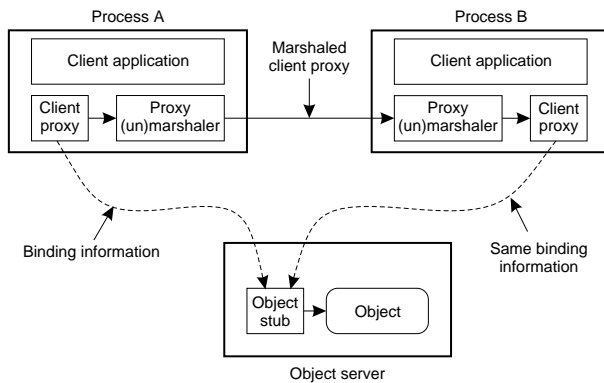
Event communication: Similar to CORBA's push-style model:



Messaging: Completely analogous to CORBA messaging.

Processes: Passing Object References

Observation: Objects are referenced by means of a local interface pointer. The question is how such pointers can be passed between different machines:



Question: Where does the proxy marshaler come from? Do we always need it?

09 – 22

Distributed Object-Based Systems/9.2 Distributed COM

Naming: Monikers

Observation: DCOM can handle only objects as temporary instances of a class. To accommodate objects that can outlive their client, something else is needed.

Moniker: A hack to support real objects

- A moniker associates data (e.g., a file), with an application or program
- Monikers can be stored
- A moniker can contain a **binding protocol**, specifying how the associated program should be “launched” with respect to the data.

| | | |
|---|--------------|--|
| 1 | Client | Calls BindMoniker at moniker |
| 2 | Moniker | Lookup CLSID and tell SCM to create object |
| 3 | SCM | Loads class object |
| 4 | Class object | Creates object, returns int. pointer |
| 5 | Moniker | Instructs object to load previously stored state |
| 6 | Object | Loads its state from file |
| 7 | Moniker | Returns interface pointer of object to client |

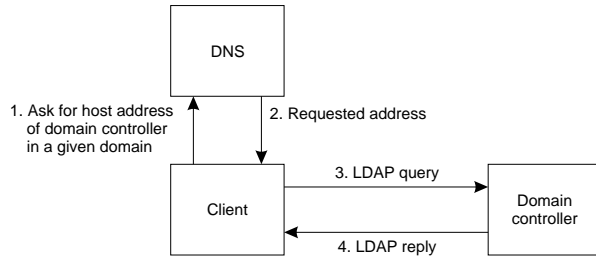
09 – 23

Distributed Object-Based Systems/9.2 Distributed COM

Active Directory

Essence: a worldwide distributed directory service, but one that does not provide location transparency.

Basics: Associate a directory service (called **domain controller**) with each domain; look up the controller using a normal DNS query:



Note: Controller is implemented as an LDAP server

Fault Tolerance

Automatic transactions: Each class object (from which objects are created), has a transaction attribute that determines how its objects behave as part of a transaction:

| Attr. value | Description |
|---------------|--|
| REQUIRES_NEW | A new transaction is always started at each invocation |
| REQUIRED | A new transaction is started if not already done so |
| SUPPORTED | Join a transaction only if caller is already part of one |
| NOT_SUPPORTED | Never join a transaction (no transaction support) |
| DISABLED | Never join a transaction, even if told to do so |

Note: Transactions are essentially executed at the level of a method invocation.

Security (1/2)

Declarative security: Register per object what the system should enforce with respect to authentication. Authentication is associated with users and user groups. There are different authentication levels:

| Auth. level | Description |
|------------------|---|
| NONE | No authentication is required |
| CONNECT | Authenticate client when first connected to server |
| CALL | Authenticate client at each invocation |
| PACKET | Authenticate all data packets |
| PACKET_INTEGRITY | Authenticate data packets and do integrity check |
| PACKET_PRIVACY | Authenticate, integrity-check, and encrypt data packets |

Security (2/2)

Delegation: A server can impersonate a client depending on a level:

| Impersonation | Description |
|---------------|--|
| ANONYMOUS | The client is completely anonymous to the server |
| IDENTIFY | The server knows the client and can do access control checks |
| IMPERSONATE | The server can invoke local objects on behalf of the client |
| DELEGATE | The server can invoke remote objects on behalf of the client |

Note: There is also support for **programmatic security** by which security levels can be set by an application, as well as the required security services (see book).


This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

-
- This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

Distributed Object-Based Systems/Globe

[illegible]

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

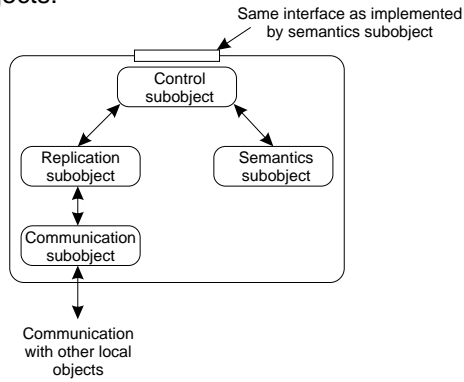


This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

[illegible]

Object Model (2/3)

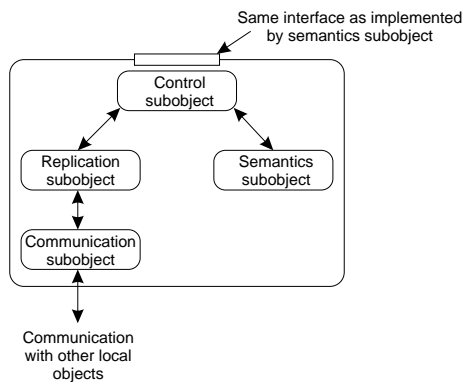
Observation: Globe attempts to separate functionality from distribution by distinguishing different local subobjects:



Semantics subobject: Contains the methods that implement the functionality of the distributed shared object

Communication subobject: Provides a (relatively simple), network-independent interface for communication between local objects

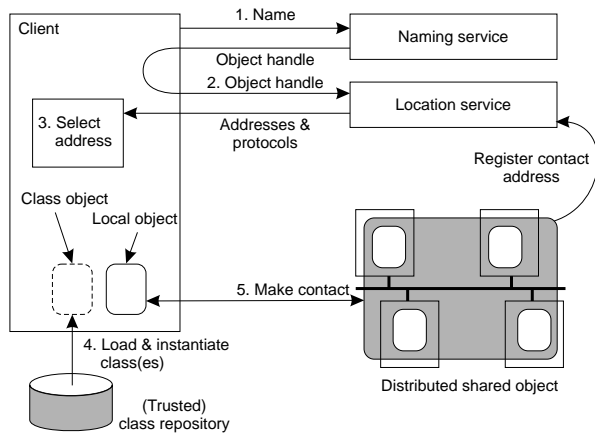
Object Model (3/3)



Replication subobject: Contains the implementation of an **object-specific** consistency protocol that controls exactly when a method on the semantics subobject may be invoked

Control subobject: Connects the user-defined interfaces of the semantics subobject to the generic, predefined interfaces of the replication subobject

Client-to-Object Binding



Observation: Globe's contact addresses correspond to CORBA's object references

Globe Services

| Service | Possible implementation | Av? |
|-----------------|--|-----|
| Collection | Separate object that holds references to other objects | No |
| Concurrency | Each object implements its own concurrency control strategy | No |
| Transaction | Separate object representing a transaction manager | No |
| Event/Notif. | Separate object per group of events (as in DCOM) | No |
| Externalization | Each object implements its own marshaling routines | Yes |
| Life cycle | Separate class objects combined with per-object implementations | Yes |
| Licensing | Implemented by each object separately | No |
| Naming | Separate service, implemented by a collection of naming objects | Yes |
| Property | Separate service, implemented by a collection of directory objects | No |
| Persistence | Implemented on a per-object basis | Yes |
| Security | Implemented per object, combined with (local) security services | Yes |
| Replication | Implemented on a per-object basis | Yes |
| Fault tolerance | Implemented per object combined with fault-tolerant servers | Yes |

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

[illegible]

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

-
-
-
-
-

[illegible]

-
-

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.



[illegible]

| Issue | CORBA | DCOM | Globe |
|----------------------|-------------------|-------------------|---------------------|
| Design goals | Interoperability | Functionality | Scalability |
| Object model | Remote objects | Remote objects | Distributed objects |
| Services | Many of its own | From environment | Few |
| Interfaces | IDL based | Binary | Binary |
| Sync. communication | Yes | Yes | Yes |
| Async. communication | Yes | Yes | No |
| Callbacks | Yes | Yes | No |
| Events | Yes | Yes | No |
| Messaging | Yes | Yes | No |
| Object server | Flexible (POA) | Hard-coded | Object dependent |
| Directory service | Yes | Yes | No |
| Trading service | Yes | No | No |
| Naming service | Yes | Yes | Yes |
| Location service | No | No | Yes |
| Object reference | Object's location | Interface pointer | True identifier |
| Synchronization | Transactions | Transactions | Only intra-object |
| Replication support | Separate server | None | Separate subobject |
| Transactions | Yes | Yes | No |
| Fault tolerance | By replication | By transactions | By replication |
| Recovery support | Yes | By transactions | No |