

High Performance Fortran (HPF)

Source:

Chapter 7 of "Designing and building parallel programs" (Ian Foster, 1995)

Question

Can't we just have a clever compiler generate a parallel program from a sequential program?

Fine-grained parallelism

```
x = a*b + c*d
```

Trivial parallelism

```
for i := 1 to 100 do
  for j := 1 to 100 do
    C[i,j] := dotproduct(A[i,*], B[:,j]);
  od
od
```

Automatic parallelism

Automatic parallelization of *any* program is extremely hard

Solutions:

- Make restrictions on source program

- Restrict kind of parallelism used

- Use semi-automatic approach

- Use application-domain oriented languages

High Performance Fortran (HPF)

Designed by a forum from industry, government, universities

Extends Fortran 90

To be used for computationally expensive numerical applications

Portable to SIMD machines, vector processors, shared-memory
MIMD and distributed-memory MIMD

Fortran 90 – Base language of HPF

Extends Fortran 77 with 'modern' features

abstract data types, modules

recursion

pointers, dynamic storage

Array operators

$A = B + C$

$A = A + 1.0$

$A(1:7) = B(1:7) + B(2:8)$

WHERE ($X \neq 0$) $X = 1.0/X$

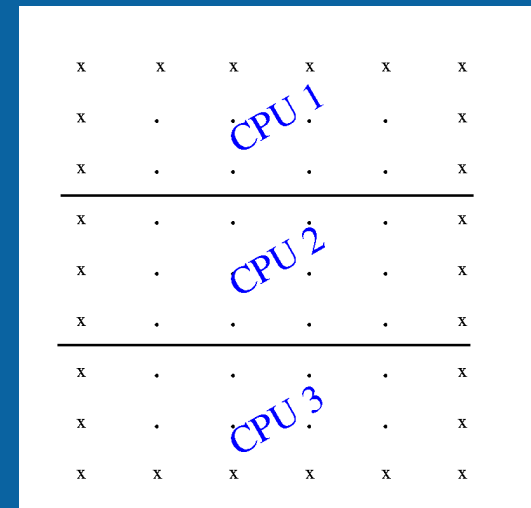
Data parallelism

Data parallelism: same operation applied to different data elements in parallel

Data parallel program: sequence of data parallel operations

Overall approach:

- Programmer does domain decomposition
- Compiler partitions operations automatically



Data may be regular (array) or irregular (tree, sparse matrix)

Most data parallel languages only deal with arrays

Data parallelism – Concurrency

Explicit parallel operations

$A = B + C$! A, B, and C are arrays

Implicit parallelism

```
do i = 1,m
  do j = 1,n
    A(i,j) = B(i,j) + C(i,j)
  enddo
enddo
```

Compiling data parallel programs

Programs are translated automatically into parallel SPMD (Single Program Multiple Data) programs

Each processor executes same program on a subset of the data

Owner computes rule:

- Each processor owns subset of the data structures
- Operations required for an element are executed by the owner
- Each processor may read (but not modify) other elements

Example

```
real s, X(100), Y(100) ! s is scalar, X and Y are arrays
```

```
X = X * 3.0           ! Multiply each X(i) by 3.0
```

```
do i = 2,99
```

```
    Y(i) = (X(i-1) + X(i+1))/2 ! Communication required
```

```
enddo
```

```
s = SUM(X)           ! Communication required
```

Arrays X and Y are distributed (partitioned)

Scalar s is replicated on each machine

HPF primitives for data distribution

Directives:

PROCESSORS: shape & size of abstract processors

ALIGN: align elements of different arrays

DISTRIBUTE: distribute (partition) an array

Directives affect performance of the program, not its result

Processors directive

```
!HPF$ PROCESSORS P(32)
```

```
!HPF$ PROCESSORS Q(4,8)
```

Mapping of abstract to physical processors not specified in HPF
(implementation-dependent)

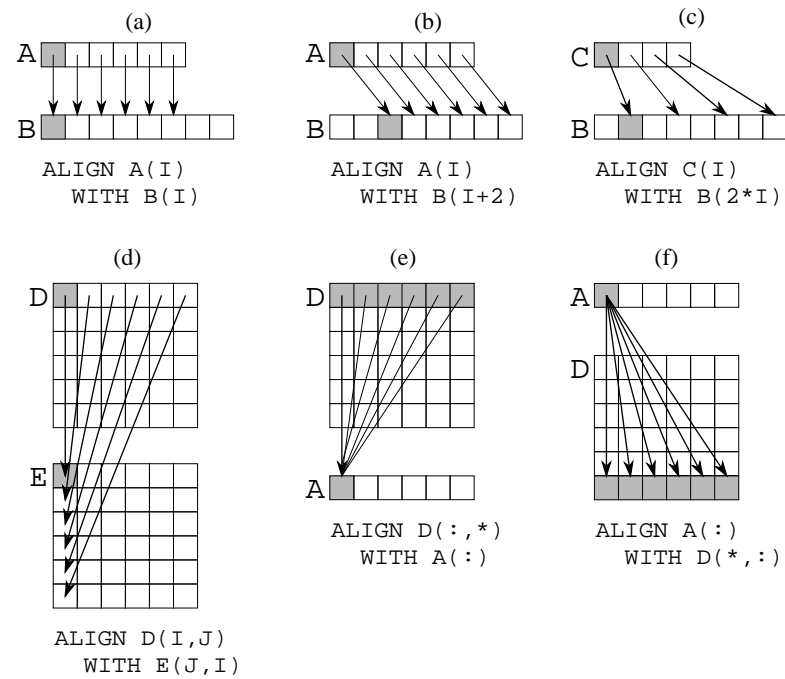
Alignment directive

Aligns an array with another array

Specifies that specific elements should be mapped to the same processor

```
real A(50), B(50), C(50,50)
!HPF$ ALIGN A(I) WITH B(I)
!HPF$ ALIGN A(I) WITH B(I+2)
!HPF$ ALIGN A(:) WITH C(*,:)
```

Figure 7.6 from Foster's book



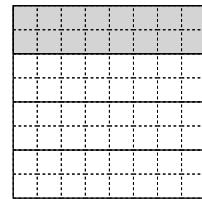
Distribution directive

Specifies how elements should be partitioned among the local memories

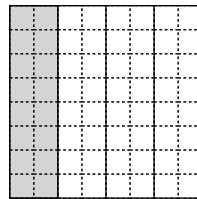
Each dimension can be distributed as follows:

<code>*</code>	<code>no distribution</code>
<code>BLOCK(n)</code>	<code>block distribution</code>
<code>CYCLIC(n)</code>	<code>cyclic distribution</code>

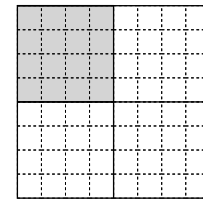
Figure 7.7 from Foster's book



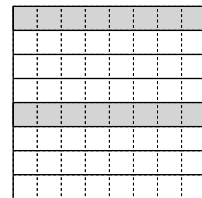
(BLOCK, *)



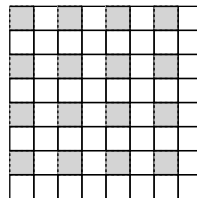
(*, BLOCK)



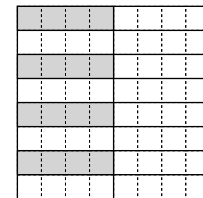
(BLOCK, BLOCK)



(CYCLIC, *)



(CYCLIC, CYCLIC)



(CYCLIC, BLOCK)

Example program

See Ian Foster, Program 7.2

```
program hpf_finite_difference
!HPF$ PROCESSORS pr(4)           ! use 4 CPUs
    real X(100, 100), New(100, 100) ! data arrays
!HPF$ ALIGN New(:, :) WITH X(:, :)
!HPF$ DISTRIBUTE X(BLOCK,*) ONTO pr ! row-wise

    New(2:99, 2:99) = (X(1:98, 2:99) + X(3:100, 2:99)
$      + X(2:99, 1:98) + X(2:99, 3:100))/4
    diffmax = MAXVAL(ABS(New-X))
end
```


Example program (2)

Use block distribution instead of row distribution

```
      program hpf_finite_difference
!HPF$ PROCESSORS pr(2,2)           ! use 2x2 grid
      real X(100, 100), New(100, 100) ! data arrays
!HPF$ ALIGN New(:, :) WITH X(:, :)
!HPF$ DISTRIBUTE X(BLOCK, BLOCK) ONTO pr ! block-wise

      New(2:99, 2:99) = (X(1:98, 2:99) + X(3:100, 2:99)
$              + X(2:99, 1:98) + X(2:99, 3:100))/4
      diffmax = MAXVAL(ABS(New-X))
      end
```

Performance

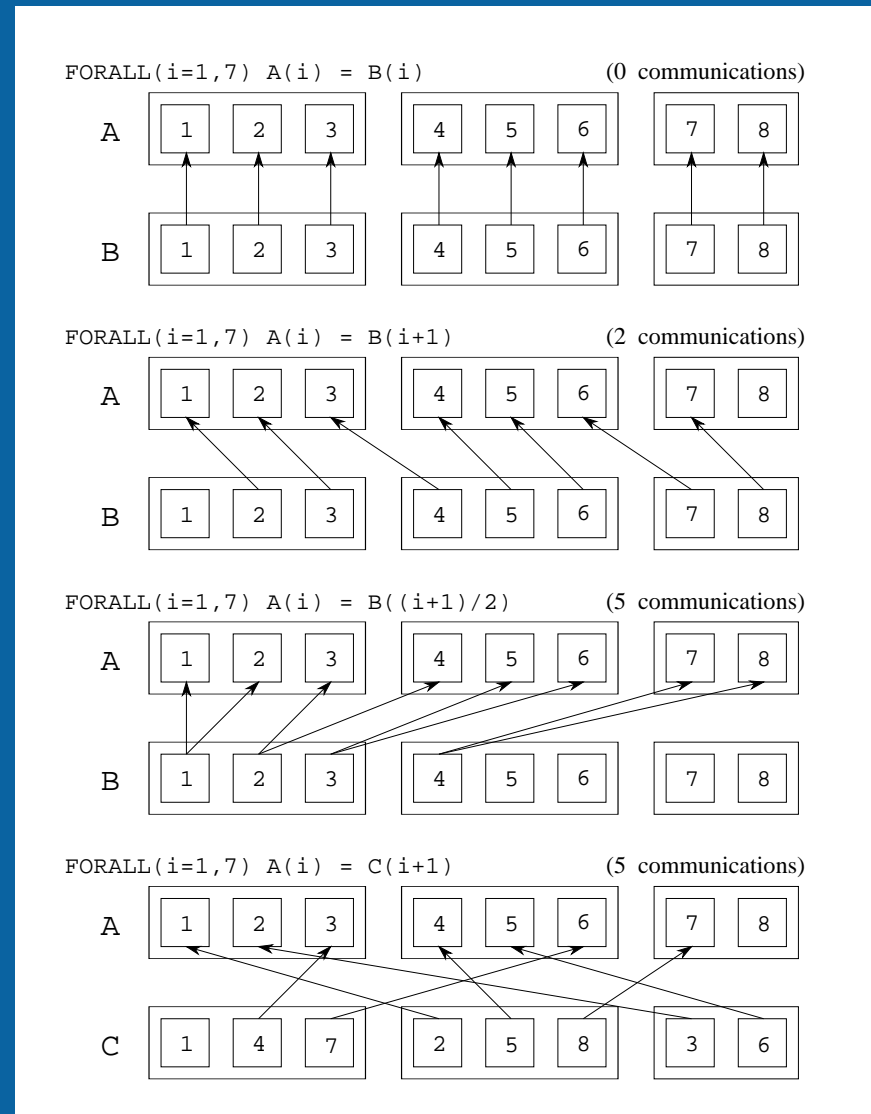
Distribution affects

- Load balance
- Amount of communication

Example (communication costs):

```
!HPF$ PROCESSORS pr(3)
      integer A(8), B(8), C(8)
!HPF$ ALIGN B(:) WITH A(:)
!HPF$ DISTRIBUTE A(BLOCK) ONTO pr
!HPF$ DISTRIBUTE C(CYCLIC) ONTO pr
```

Figure 7.9 from Foster's book



Conclusions

High-level model

- User specifies data distribution

- Compiler generates parallel program + communication

More restrictive than general message passing model (only data parallelism)

Restricted to array-based data structures

HPF programs will be easy to modify, which enhances portability

Changing data distribution only requires changing directives