

Introduction to Parallel Programming

- Language notation: message passing
- 5 parallel algorithms of increasing complexity:
 - ★ Matrix multiplication
 - ★ Successive overrelaxation
 - ★ All-pairs shortest paths
 - ★ Linear equations
 - ★ Search problem

Message Passing

- SEND(destination, message)
 - ★ blocking: wait until message has arrived
 - ★ nonblocking: continue immediately
- RECEIVE(source, message)
- RECEIVE-FROM-ANY(message)
 - ★ blocking: wait until message is available
 - ★ nonblocking: test if message is available

Parallel Matrix Multiplication

- Given two $N \times N$ matrices A and B
- Compute $C = A \times B$
- $C_{ij} = A_{i1}B_{1j} + A_{i2}B_{2j} + .. + A_{iN}B_{Nj}$

$$\begin{pmatrix} \text{XXXXX} \\ \text{XXXXX} \\ \text{XXXXX} \\ \text{XXXXX} \end{pmatrix} \mathbf{X} \begin{pmatrix} \text{XXXXX} \\ \text{XXXXX} \\ \text{XXXXX} \\ \text{XXXXX} \end{pmatrix} = \begin{pmatrix} \text{XXXXX} \\ \text{XXXXX} \\ \text{XXXXX} \\ \text{XXXXX} \end{pmatrix}$$

Sequential Matrix Multiplication

```
for (i = 1; i <= N; i++)  
    for (j = 1; j <= N; j++)  
        C[i,j] = 0;  
        for (k = 1; k <= N; k++)  
            C[i,j] += A[i,k] * B[k,j];
```

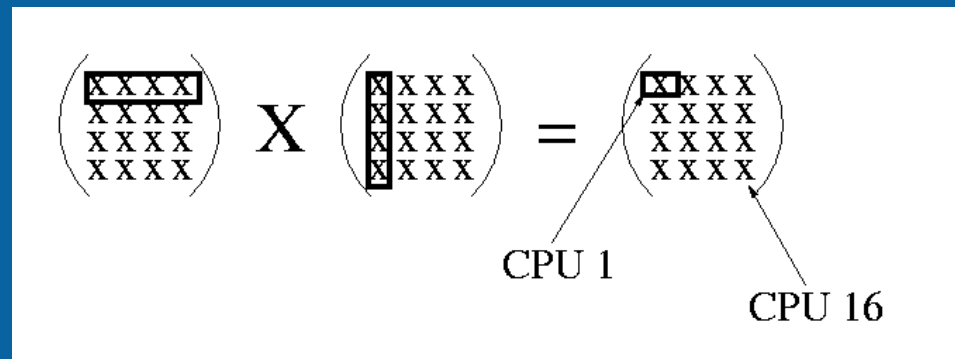
- The order of the operations is overspecified
- Everything can be computed in parallel

Parallel Algorithm 1

Each processor computes 1 element of C

Requires N^2 processors

Need 1 row of A and 1 column of B as input



Parallel Algorithm 1

Master (processor 0):

```
for (i = 1; i <= N; i++)
    for (j = 1; j <= N; j++)
        SEND(p++, A[i,*], B[*],j], i, j);
for (x = 1; x <= N*N; x++)
    RECEIVE_FROM_ANY(&result, &i, &j);
    C[i,j] = result;
```

Slaves:

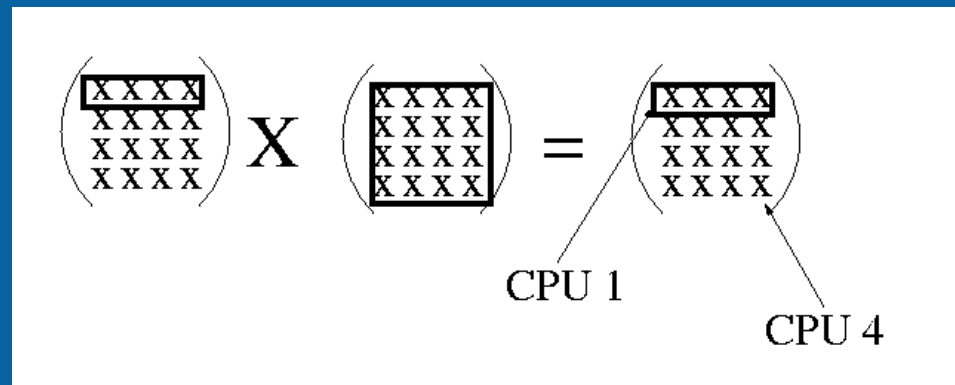
```
int Aix[N], Bxj[N], Cij;
RECEIVE(0, &Aix, &Bxj, &i, &j);
Cij = 0;
for (k = 1; k <= N; k++) Cij += Aix[k] * Bxj[k];
SEND(0, Cij , i, j);
```

Parallel Algorithm 2

Each processor computes 1 row (N elements) of C

Requires N processors

Need entire B matrix and 1 row of A as input



Parallel Algorithm 2

Master (processor 0):

```
for (i = 1; i <= N; i++)  
    SEND(i, A[i,*], B[*,*], i);  
for (x = 1; x <= N; x++)  
    RECEIVE_FROM_ANY(&result, &i);  
    C[i,*] = result[*];
```

Slaves:

```
int Aix[N], B[N,N], C[N];  
RECEIVE(0, &Aix, &B, &i);  
for (j = 1; j <= N; j++)  
    C[j] = 0;  
    for (k = 1; k <= N; k++) C[j] += Aix[k] * B[j,k];  
SEND(0, C[*] , i);
```


Problem: need larger granularity

So far, each parallel task needs as much communication as computation

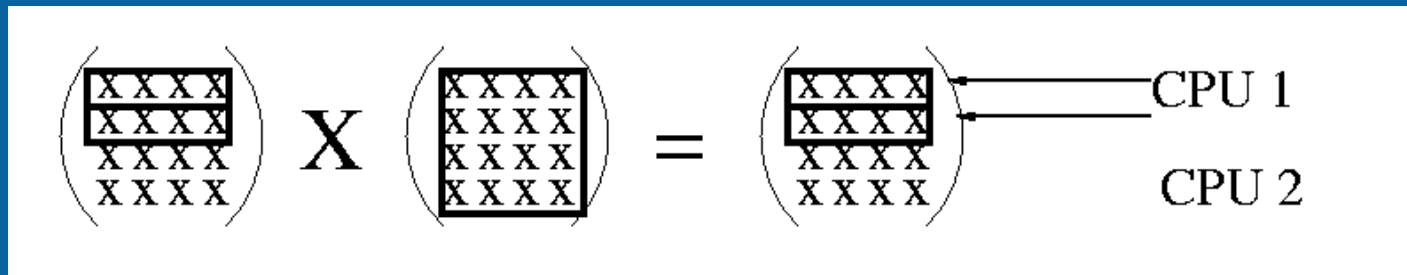
Assumption: $N \gg P$ (i.e. we solve a *large* problem)

Assign many rows to each processor

Parallel Algorithm 3

Each processor computes $\frac{N}{P}$ rows of C

Need entire B matrix and $\frac{N}{P}$ rows of A as input



Parallel Algorithm 3

Master (processor 0):

```
int result[N, N/nprocs];
int inc = N / nprocs;    /* number of rows per cpu */
int lb = 1;
for (i = 1; i <= nprocs; i++)
    SEND(i, A[lb .. lb+inc-1, *], B[*,*], lb, lb+inc-1);
    lb += inc;
for (x = 1; x <= nprocs; x++)
    RECEIVE_FROM_ANY(&result, &lb);
    for (i = 1; i <= N/nprocs; i++)
        C[lb+i-1, *] = result[i, *];
```

Parallel Algorithm 3 (Cnt'd)

Slaves:

```
int A[N/nprocs, N], B[N,N], C[N/nprocs, N];  
RECEIVE(0, &A, &B, &lb, &ub);  
for (i = lb; i <= ub; i++)  
    for (j = 1; j <= N; j++)  
        C[i,j] = 0;  
        for (k = 1; k <= N; k++)  
            C[i,j] += A[i,k] * B[k,j];  
SEND(0, C[*,*], lb);
```

Comparison

alg.	parallelism (#jobs)	communication per job	comput. per job	ratio comp/comm
1.	N^2	$N + N + 1$	N	$O(1)$
2.	N	$N + N^2 + N$	N^2	$O(1)$
3.	P	$\frac{N^2}{P} + N^2 + \frac{N^2}{P}$	$\frac{N^3}{P}$	$O(\frac{N}{P})$

If $N \gg P$, algorithm 3 will have low communication overhead

Its *grain size* is high

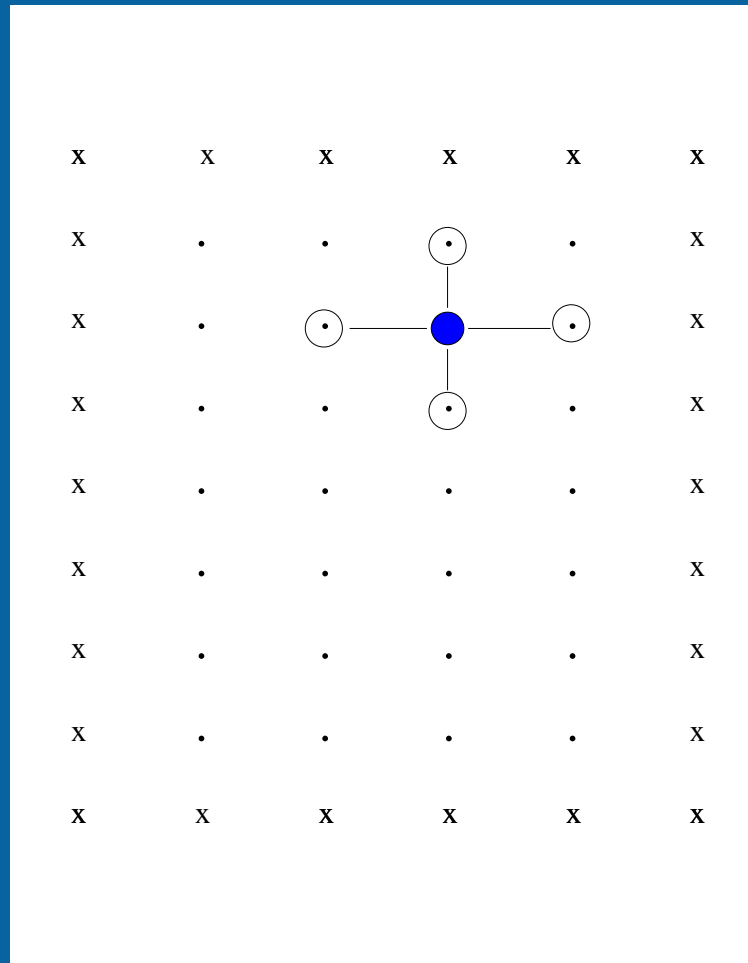
Discussion

- Matrix multiplication is trivial to parallelize
- Getting good performance is a problem
- Need right grain size
- Need large input problem

SOR example

X	X	X	X	X	X
X	X
X	X
X	X
X	X
X	X
X	X
X	X
X	X	X	X	X	X

SOR example



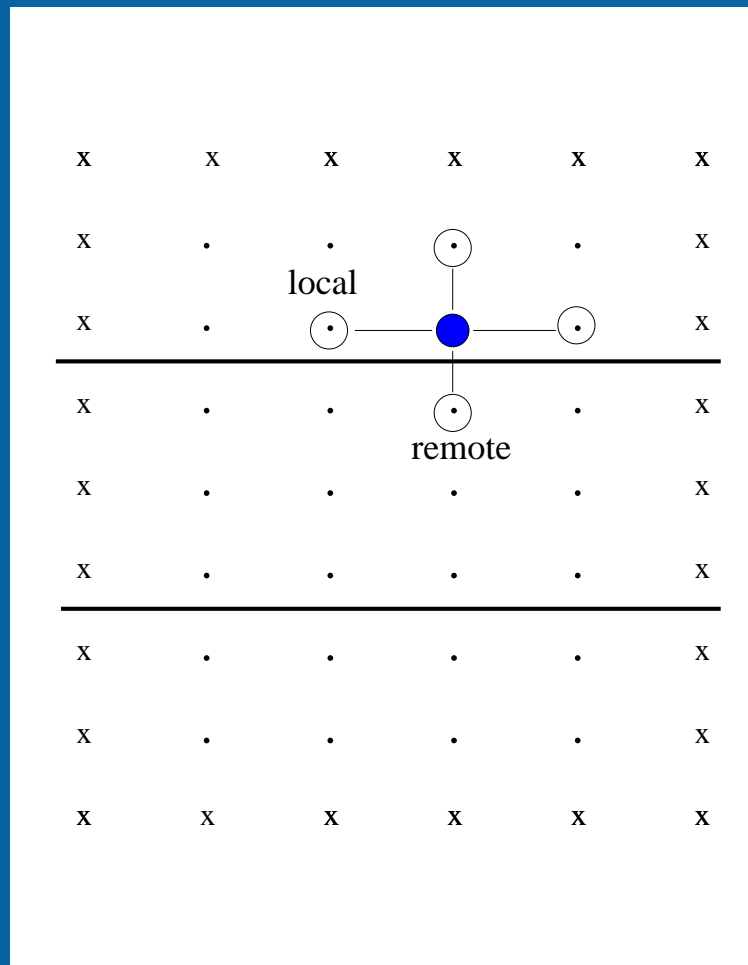
Parallelizing SOR

- Domain decomposition on the grid
- Each processor owns N/P rows
- Need communication between neighbors to exchange elements at processor boundaries

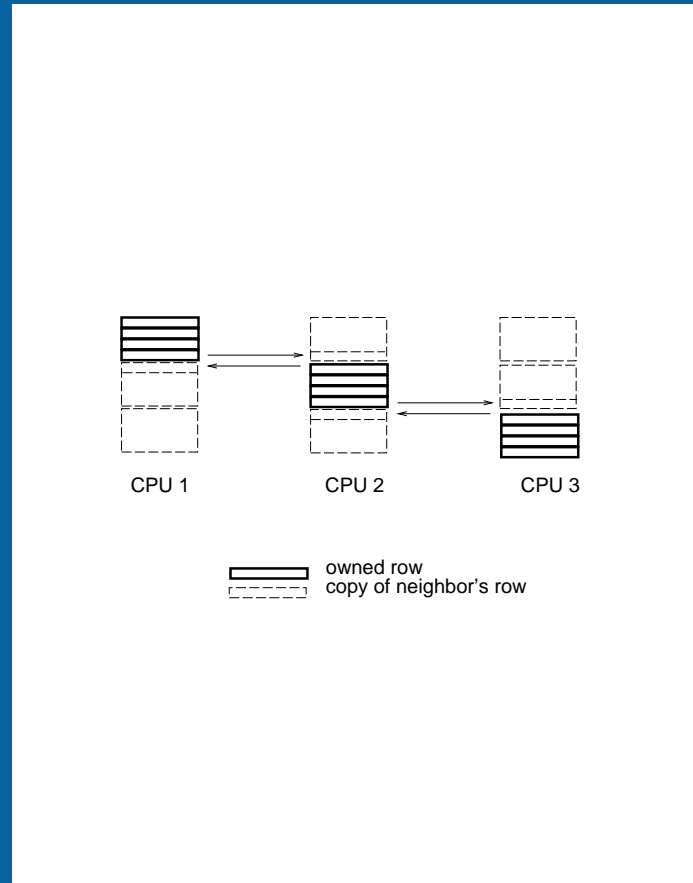
SOR example partitioning

X	X	X	X	X	X
X	X
X	X
<hr/>					
X	X
X	X
X	X
<hr/>					
X	X
X	X
X	X	X	X	X	X

SOR example partitioning



Communication scheme



Each CPU communicates with left & right neighbor (if existing)

Performance of SOR

Communication and computation during each iteration:

- Each processor sends/receives 2 messages with M reals
- Each processor computes $N/P * M$ updates

The algorithm will have good performance if

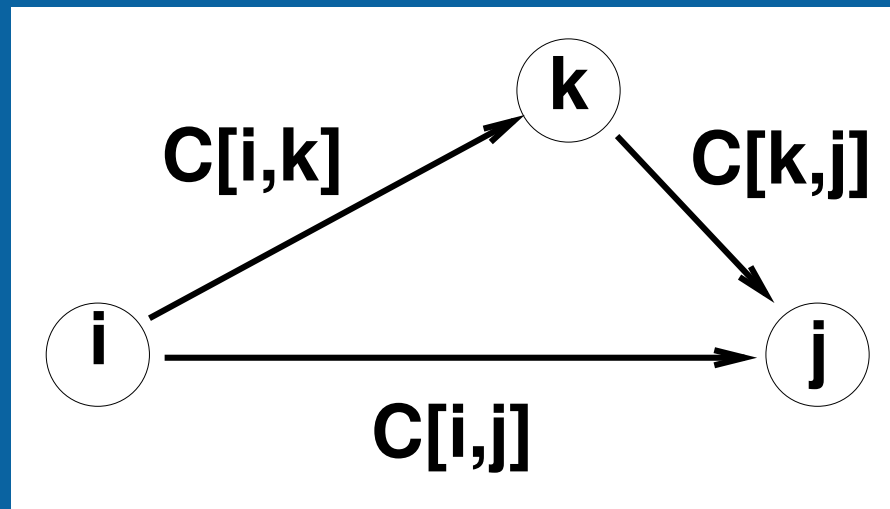
- Problem size is large: $N \gg P$
- Message exchanges can be done in parallel

All-pairs Shorts Paths (ASP)

- Given a graph G with a distance table C :
 - ★ $C[i,j]$ = length of direct path from node i to node j
- Compute length of shortest path between any two nodes in G

Floyd's Sequential Algorithm

- Basic step:



```
for (k = 1; k <= N; k++)  
  for (i = 1; i <= N; i++)  
    for (j = 1; j <= N; j++)  
      C[i,j] = MIN(C[i,j], C[i,k] + C[k,j])
```

Parallelizing ASP

- Distribute rows of C among the P processors
- During iteration k , each processor executes
$$C[i,j] = \text{MIN}(C[i,j], C[i,k] + C[k,j]);$$
on its own rows i , so it needs these rows and row k
- Before iteration k , the processor owning row k sends it to all the others

Parallel ASP Algorithm

```
int lb, ub;          /* lower/upper bound for this CPU */
int rowK[N], C[lb:ub, N]; /* pivot row ; matrix */

for (k = 1; k <= N; k++)
    if (k >= lb && k <= ub) /* do I have it? */
        rowK = C[k,*];
        for (p = 1; p <= nproc; p++) /* broadcast row */
            if (p != myprocid) SEND(p, rowK);
    else
        RECEIVE_FROM_ANY(&rowK); /* receive row */
    for (i = lb; i <= ub; i++) /* update my rows */
        for (j = 1; j <= N; j++)
            C[i,j] = MIN(C[i,j], C[i,k] + rowK[j]);
```

Parallel ASP Algorithm

```
int lb, ub;          /* lower/upper bound for this CPU */
int rowK[N], C[lb:ub, N]; /* pivot row ; matrix */

for (k = 1; k <= N; k++)

    for (i = lb; i <= ub; i++) /* update my rows */
        for (j = 1; j <= N; j++)
            C[i,j] = MIN(C[i,j], C[i,k] + rowK[j]);
```

Parallel ASP Algorithm

```
int lb, ub;          /* lower/upper bound for this CPU */
int rowK[N], C[lb:ub, N]; /* pivot row ; matrix */

for (k = 1; k <= N; k++)
    if (k >= lb && k <= ub) /* do I have it? */
        rowK = C[k,*];
        for (p = 1; p <= nproc; p++) /* broadcast row */
            if (p != myprocid) SEND(p, rowK);
    else
        RECEIVE_FROM_ANY(&rowK); /* receive row */
    for (i = lb; i <= ub; i++) /* update my rows */
        for (j = 1; j <= N; j++)
            C[i,j] = MIN(C[i,j], C[i,k] + rowK[j]);
```

Performance Analysis ASP

Per iteration:

- 1 CPU sends $P - 1$ messages with N integers
- Each CPU does $\frac{N}{P} \times N$ comparisons

Communication/computation ratio is small if $N \gg P$

... but, is the Algorithm Correct?

?

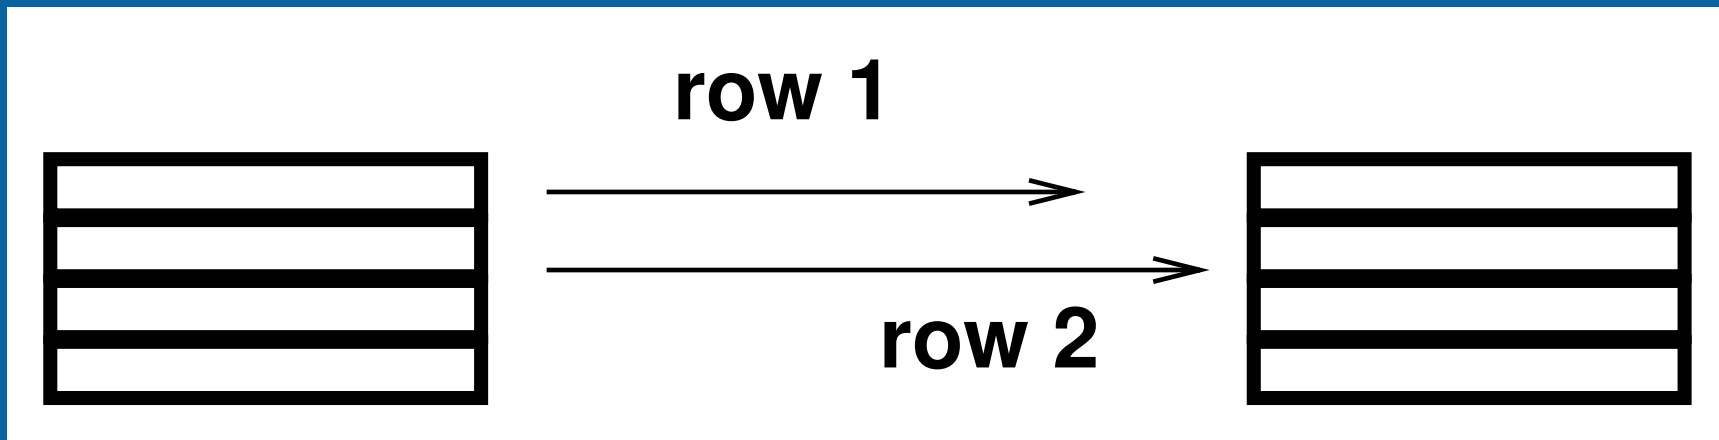
Parallel ASP Algorithm

```
int lb, ub;
int rowK[N], C[lb:ub, N];

for (k = 1; k <= N; k++)
    if (k >= lb && k <= ub)
        rowK = C[k,*];
        for (p = 1; p <= nproc; p++)
            if (p != myprocid) SEND(p, rowK);
    else
        RECEIVE_FROM_ANY(&rowK);
    for (i = lb; i <= ub; i++)
        for (j = 1; j <= N; j++)
            C[i,j] = MIN(C[i,j], C[i,k] + rowK[j]);
```

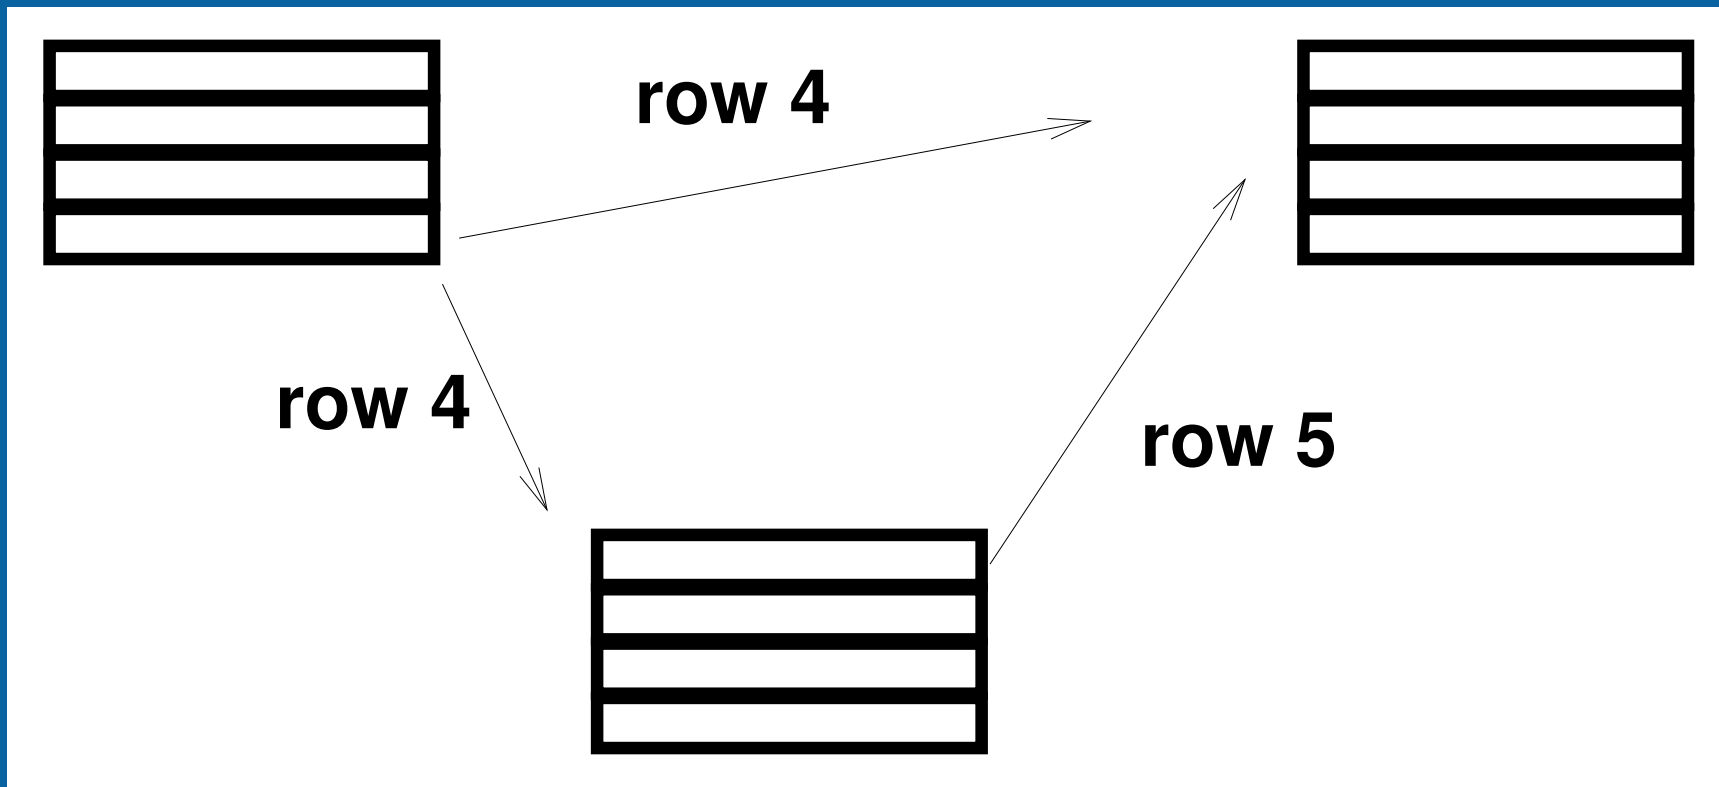

Non-FIFO Message Ordering

Row 2 may be received before row 1



FIFO Ordering

Row 5 may be received before row 4



Correctness

Problems:

- Asynchronous non-FIFO SEND
- Messages from different senders may overtake each other

Correctness

Problems:

- Asynchronous non-FIFO SEND
- Messages from different senders may overtake each other

Solutions:

Correctness

Problems:

- Asynchronous non-FIFO SEND
- Messages from different senders may overtake each other

Solutions:

- Synchronous SEND (less efficient)

Correctness

Problems:

- Asynchronous non-FIFO SEND
- Messages from different senders may overtake each other

Solutions:

- Synchronous SEND (less efficient)
- Barrier at the end of outer loop (extra communication)

Correctness

Problems:

- Asynchronous non-FIFO SEND
- Messages from different senders may overtake each other

Solutions:

- Synchronous SEND (less efficient)
- Barrier at the end of outer loop (extra communication)
- Order incoming messages (requires buffering)

Correctness

Problems:

- Asynchronous non-FIFO SEND
- Messages from different senders may overtake each other

Solutions:

- Synchronous SEND (less efficient)
- Barrier at the end of outer loop (extra communication)
- Order incoming messages (requires buffering)
- RECEIVE(cpu, msg) (more complicated)

Linear equations

- Linear equations:

$$a_{1,1}x_1 + a_{1,2}x_2 + \dots a_{1,n}x_n = b_1$$

...

$$a_{n,1}x_1 + a_{n,2}x_2 + \dots a_{n,n}x_n = b_n$$

- Matrix notation: $Ax = b$
- Problem: compute x , given A and b
- Linear equations have many important applications
Practical applications need huge sets of equations

Solving a linear equation

- Two phases:

Upper-triangularization $\rightarrow Ux = y$

Back-substitution $\rightarrow x$

- Most computation time is in upper-triangularization

- Upper-triangular matrix:

$$U[i,i] = 1$$

$$U[i,j] = 0 \text{ if } i > j$$

Sequential Gaussian elimination

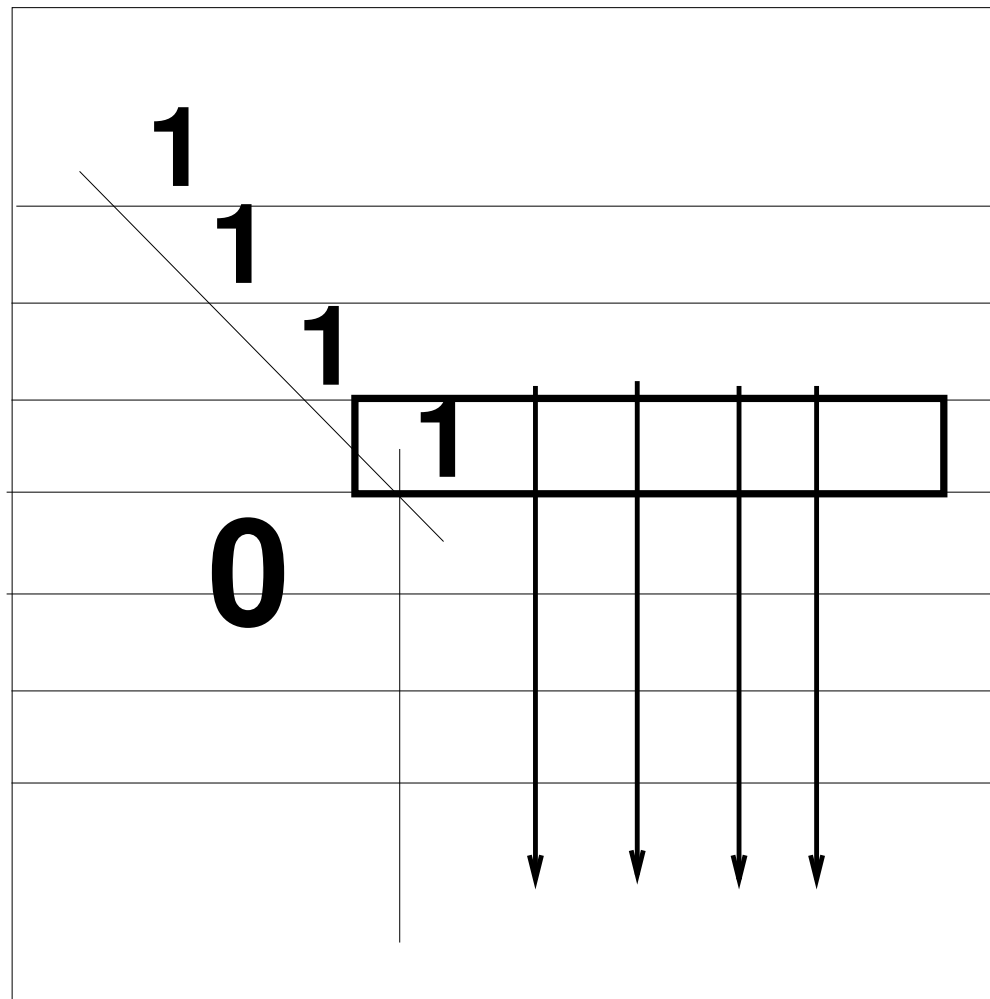
```
for (k = 1; k <= N; k++)  
    for (j = k+1; j <= N; j++)  
        A[k,j] = A[k,j] / A[k,k]  
y[k] = b[k] / A[k,k]  
A[k,k] = 1  
for (i = k+1; i <= N; i++)  
    for (j = k+1; j <= N; j++)  
        A[i,j] = A[i,j] - A[i,k] * A[k,j]  
    b[i] = b[i] - A[i,k] * y[k]  
    A[i,k] = 0
```

- Converts $Ax = b$ into $Ux = y$
- Sequential algorithm uses $\frac{2}{3}N^3$ operations

Parallelizing Gaussian elimination

- Row-wise partitioning scheme
 - ★ Each CPU gets one row (*striping*)
 - ★ Execute one (outer-loop) iteration at a time
- Communication requirement:
 - ★ During iteration k , CPUs $P_{k+1} \dots P_{n-1}$ need part of row k
 - ★ This row is stored on CPU P_k
 - ★ \rightarrow need partial broadcast (multicast)

Communication

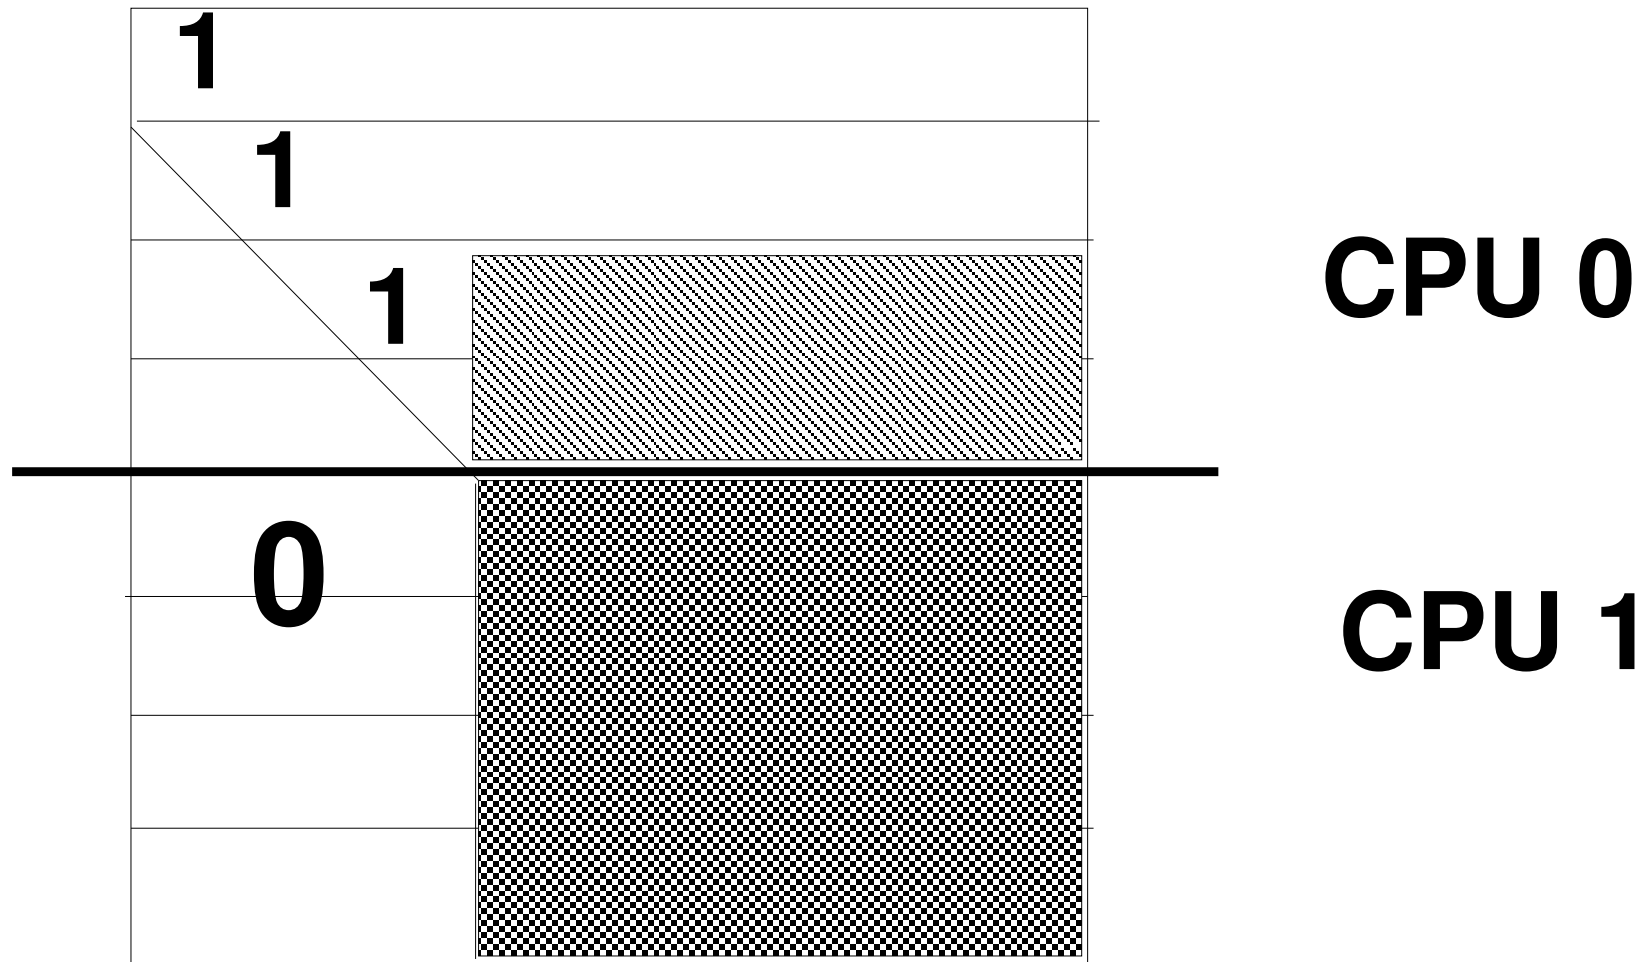


multicast

Performance problems

- Communication overhead (multicast)
- Load imbalance
 - CPUS $P_0...P_k$ are idle during iteration k
- In general, number of CPUS is less than n
 - Choice between block-striped and cyclic-striped distribution
- Block-striped distribution has high load-imbalance
- Cyclic-striped distribution has less load-imbalance

Block-striped distribution



A Search Problem

Given an array $A[1..N]$ and an item x , check if x is present in A

```
int present = false;  
for (i = 1; !present && i <= N; i++)  
    if (A[i] == x) present = true;
```

Parallel Search on 2 CPUs

```
int lb, ub;
int A[lb:ub];

for (i = lb; i <= ub; i++)
    if (A[i] == x)
        print("Found item");
        SEND(1-cpuid); /* send other CPU empty message*/
        exit();
/* check message from other CPU: */
if (NONBLOCKING_RECEIVE(1-cpuid)) exit()
```

Performance Analysis

How much faster is the parallel program than the sequential program for $N=100$?

Performance Analysis

How much faster is the parallel program than the sequential program for $N=100$?

1. if x not present \Rightarrow factor 2

Performance Analysis

How much faster is the parallel program than the sequential program for $N=100$?

1. if x not present \Rightarrow factor 2
2. if x present in $A[1 \dots 50] \Rightarrow$ factor 1

Performance Analysis

How much faster is the parallel program than the sequential program for $N=100$?

1. if x not present \Rightarrow factor 2
2. if x present in $A[1 \dots 50] \Rightarrow$ factor 1
3. if $A[51] = x \Rightarrow$ factor 51

Performance Analysis

How much faster is the parallel program than the sequential program for $N=100$?

1. if x not present \Rightarrow factor 2
2. if x present in $A[1 \dots 50] \Rightarrow$ factor 1
3. if $A[51] = x \Rightarrow$ factor 51
4. if $A[75] = x \Rightarrow$ factor 3

Performance Analysis

How much faster is the parallel program than the sequential program for $N=100$?

1. if x not present \Rightarrow factor 2
2. if x present in $A[1 \dots 50] \Rightarrow$ factor 1
3. if $A[51] = x \Rightarrow$ factor 51
4. if $A[75] = x \Rightarrow$ factor 3

In case 2 the parallel program does more work than the sequential program \Rightarrow *search overhead*

Performance Analysis

How much faster is the parallel program than the sequential program for $N=100$?

1. if x not present \Rightarrow factor 2
2. if x present in $A[1 \dots 50] \Rightarrow$ factor 1
3. if $A[51] = x \Rightarrow$ factor 51
4. if $A[75] = x \Rightarrow$ factor 3

In case 2 the parallel program does more work than the sequential program \Rightarrow *search overhead*

In cases 3 and 4 the parallel program does less work \Rightarrow *negative search overhead*

Discussion

Several kinds of performance overhead

Discussion

Several kinds of performance overhead

- Communication overhead

Discussion

Several kinds of performance overhead

- Communication overhead
- Load imbalance

Discussion

Several kinds of performance overhead

- Communication overhead
- Load imbalance
- Search overhead

Discussion

Several kinds of performance overhead

- Communication overhead
- Load imbalance
- Search overhead

Making algorithms correct is nontrivial

Discussion

Several kinds of performance overhead

- Communication overhead
- Load imbalance
- Search overhead

Making algorithms correct is nontrivial

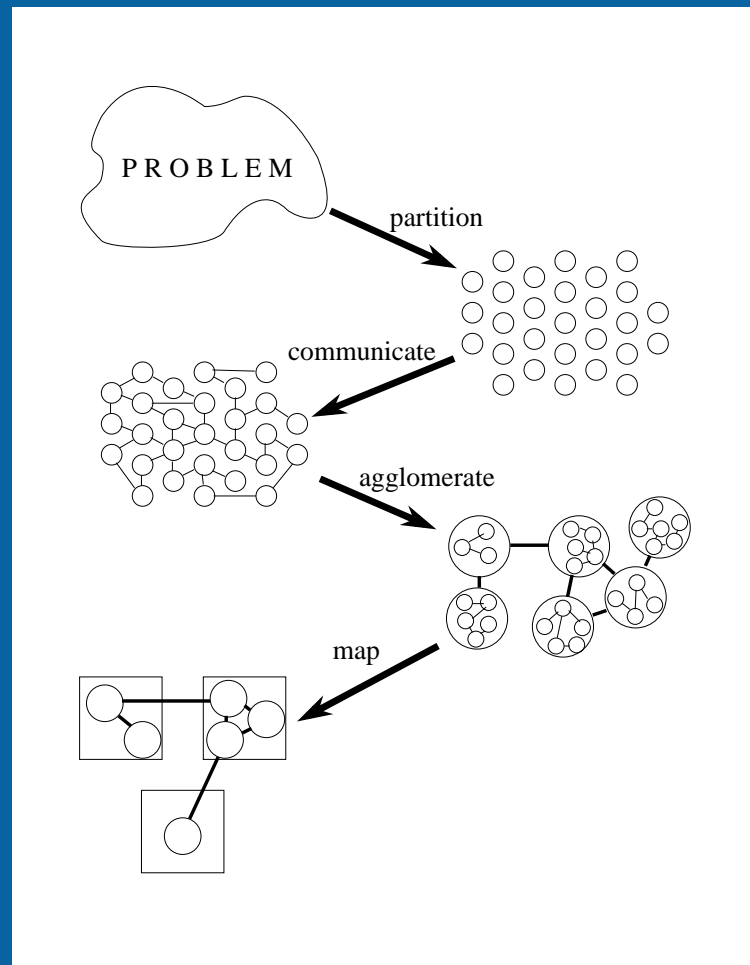
- Message ordering

Designing Parallel Algorithms

Source: Designing and building parallel programs (Ian Foster, 1995)

- Partitioning
- Communication
- Agglomeration
- Mapping

Figure 2.1 from Foster's book



Partitioning

- Domain decomposition
 - ★ Partition the data
 - ★ Partition computations on data (owner-computes rule)
- Functional decomposition
 - ★ Divide computations into subtasks
 - ★ E.g. search algorithms

Communication

- Analyze data-dependencies between partitions
- Use communication to transfer data
- Many forms of communication, e.g.
 - ★ Local communication with neighbors (SOR)
 - ★ Global communication with all processors (ASP)
 - ★ Synchronous (blocking) communication
 - ★ Asynchronous (nonblocking) communication

Agglomeration

Reduce communication overhead by

- increasing granularity
- improving locality

Mapping

- On which processor to execute each subtask?
- Put concurrent tasks on different CPUs
- Put frequently communicating tasks on same CPU?
- Avoid load imbalances

Summary

Hardware and software models

Example applications

- Matrix multiplication
- Successive overrelaxation
- All-pairs shortest paths
- Linear equations
- Search problem

Designing parallel algorithms