

Performance Metrics

$$\text{Speedup} = \frac{\text{Time on 1 CPU}}{\text{Time on } P \text{ CPUs}}$$

$$\text{Efficiency} = \frac{\text{speedup}}{p}$$

What do we compare against?

parallel program on 1 CPU?

equivalent sequential program?

best sequential program?

Example: sorting

Suppose a parallel bubble sort program obtains 100% efficiency

Efficiency relative to quicksort: $\frac{\frac{N \log N}{N^2/p}}{p} = \frac{\log N}{N}$

How to Cheat with Speedups (1)

Program with hardware floating point support

Sequential program: 100 sec

Parallel program on 10 CPUs:

 computation: 10 sec

 communication: 10 sec

Speedup = $100 / (10 + 10) = 5$

How to Cheat with Speedups (2)

Without hardware floating point support

Sequential program: 400 sec

Parallel program on 10 CPUs:

 computation: 40 sec

 communication: 10 sec

Speedup = $400 / (40 + 10) = 8$

Can Speedup be Superlinear?

Superlinear speedup: $\text{speedup} \geq p$

Negative search overhead (for search problems)

Applies to certain input problems, not to average case

Sequential program could simulate parallel program to get same behavior

Better caching behavior on parallel systems

Amdahl's Law

Let f = fraction of code that must be performed sequentially

$$\text{Speedup} \leq \frac{1}{f + (1-f)/p}$$

Example: $f=10\% \Rightarrow \text{speedup} \leq 10$, even if p is infinite!

So much for parallel computing?

Amdahl's Law – the True Story

In practice, f depends on the problem size

Goal of parallel computing is to solve *large* problems fast

If problem size increases, maximum speedup increases

Scaled speedup =
$$\frac{\text{time sequential program would have taken}}{\text{time parallel program}}$$

How to Measure Performance

Use wall clock time

```
int start, stop;  
start = get_time();  
do computations ...  
stop = get_time();  
print(stop-start);
```

Exclude initialization

Measures file I/O performance

Avoid debugging statements

Performance Debugging

Try to understand performance behavior

Just attributing poor performance to 'communication overhead' is unacceptable

Determine where the time is spent

Measure number of messages per second, data volume per second, idle time